
20% if you write “**I don’t know**” for any (sub)question, with a cap of 6.

“**I don’t know**” can be written for subquestions that are marked with “partial IDK”.

Avoid detailed pseudocode; instead, explain in words.

Common algorithms like binary search, sorting, graph algorithms done in course can be used directly.

Only 1 page is allowed for every question from Q1–Q6, except Q5(d) which is allowed 1 page of its own.

Q1 [4+4=8 points (partial IDK)] Consider a data structure that maintains a forest of rooted trees, and further supports three operations.

1. *MakeTree*(x): Given an element x , create a tree whose only node is x .
2. *Depth*(x): Given a node x , return the depth of x in the tree containing x (depth of the root node of any tree is defined to be 0).
3. *Join*(r, x): Given a root node r , and another node x not in the same tree as that of r (x need not be a root), make r the child of x .

Suppose that we use reversed trees to implement this data structure (similar to what we did for disjoint-sets): we use $v.p$ to denote the parent of a node v , and set $v.p = v$ if v is a root node. Suppose further that we implement *Join*(r, x) by setting $r.p = x$ and *Depth*(x) by following the path from x up to a root, returning a count of all nodes on the path other than x . Clearly, the complexities of *MakeTree* and *Join* are $O(1)$ and that of *Depth* is linear in the depth of the given node.

Show that the worst-case total running time of any sequence of m operations (consisting of a mix of *MakeTree*(), *Depth*(), *Join*()) (a) is $O(m^2)$ and (b) is $\Omega(m^2)$. Think carefully what each of the bounds mean.

Q2 [1+1+1+1=4 points] Fill in the blanks below to complete a recursive algorithm that identifies Hamiltonian cycle in an undirected graph G (or returns NULL if none exists). The function is initially called as *FindHC*(G). The algorithm may make calls to another algorithm *HasHC*(H) that determines if H has a Hamiltonian cycle. Assuming that *HasHC* runs in polynomial time, *FindHC* should also run in polynomial time. Fill in only the steps (a)–(d) below; write “no-op” if nothing should be done for any of the steps. *FindHC*(G) should return a *sequence* of vertices on any Hamiltonian cycle (e.g., for a cycle $a - b - c - d - a$, the function can return any of $[a, b, c, d]$, $[b, c, d, a]$, $[c, d, a, b]$, $[d, a, b, c]$); however, *only one* sequence should be returned. Feel free to use known polynomial-time graph algorithms, and avoid detailed pseudocode/Java/C++/etc. You cannot use any global variable or modify the input/output signature of the function. You may add a 2-3 line explanation for your steps.

```
def FindHC(graph H): // should return NULL if no cycle exists

    // (a) fill the base case

    // (b) add anything you want to perform before the for-loop

    for edge e in H:
        // there is a Hamiltonian cycle in H after removing the edge e
        if HasHC(H-{e}) returns true:
            // (c) fill what to do here

    else:
        no-op // I am not kiddin'

    // (d) anything you want to perform after the loop, including a return statement
```

Q3 [5 points] Show that for any integer r , it is not possible to construct a polynomial time r -absolute-approximation algorithm for 0-1 Knapsack with integer values and integer weights (unless $P \neq NP$).


Q4 [1+3+1=5 points] Suppose you are given a list L of n distinct points in a two-dimensional plane; there could be points that have common X coordinates or common Y coordinates, but not both, and further, the points in L are in the order of increasing X coordinates.

A point $p \in L$ is defined to be maximal in L if no other point in L is **both** above p and to the right of p . In the example below, L has 5 maximal points.

```
def CountMaximal(list L of points sorted in X coord):
    // (a) add base case
```

```
    LL = left half of points in L
    LR = right half of points in L
```


```
    // (b) finish the rest, must call CountMaximal(LL) and CountMaximal(LR)
```



Complete the divide-and-conquer algorithm above that returns the set of maximal points in L in $O(n \log n)$ time; if you require, you can design **CountMaximal** to compute and return other information as well. The sub-problems for the divide step has been computed for you; fill in the rest. You may add 2-3 lines explaining your idea. (c) You should explain the running time with the help of a recurrence relation (no need to solve the recurrence).

Q5 [1+3+1+7=12 (partial IDK)] A “tee” is a subgraph that consists of three equal-length paths connected to one node (all the three paths should have the same number of edges).

The weight of a tee is the sum of the weights of its edges. The figure shows a graph on the left with edge weights 1 on all edges and a tee on the right with weight 12. We want to find a tee in a weighted undirected graph with the largest weight – call this the **FINDTEE** problem.



(a) Show how to define a decision problem named **DECTEE** corresponding to **FINDTEE**. (b) Explain how a polynomial algorithm for **DECTEE** can be used to solve **FINDTEE** in polynomial time; explain both algorithm and complexity analysis in words. (c) Show that **DECTREE** is an NP problem by describing a verifier. (d) Prove that **DECTREE** is NP-complete (hint: $stHamPath(G, s, t)$, that decides if there is a Hamiltonian path from s to t in a graph G , is NP-complete).

Q6 [6 points] A computer has to pick one of two tasks, Task-0 and Task-1, for each of the next n time slots. A table $U[0, 1][1 \dots n]$ is available in which $U[0][i]$ indicates the amount of utility that will be generated if Task-0 is scheduled for the i th slot, and $U[1][i]$ indicates the utility generated by scheduling Task-1 for the i -slot. However, the computer also has to pay a preemption penalty of S utility everytime one task is preempted for another. The computer has to figure out which task should be scheduled in which slot so that the total utility after the n th slot is maximized.

	$n = 3, S = 2$		
$U[0] =$	3	1	2
$U[1] =$	1	4	2

In this example, the schedule (Task-0, Task-0, Task-1) has utility $3 + 1 + 2 = 6$ and (Task-0, Task-1, Task-0) has utility $3 - 2 + 4 - 2 + 2 = 5$ due to switching penalties.

Design either a dynamic programming algorithm or a graph reduction algorithm to solve this optimization problem given the table U as input (**use only one of the approaches**).

For a DP-based approach, (**a [1 point]**) specify the English description of a function, (**b [3 points]**) a recursive formula (including a base case) to compute that function, (**c [1 point]**) describe how to compute the optimal utility from the function values, and (**d [1 point]**) explain the time complexity for doing that. Add 2-3 lines explaining the correctness of your formula in (b).

For a graph reduction based approach, explain how to construct a graph G on which we can run an algorithm for longest (with the largest weight) path on G between **one pair of vertices** to solve the scheduling question. Describe (**a [2 point]**) what do the vertices, edges and their weights represent with respect to the scheduling problem, (**b [1 point]**) the number of vertices and edges of G (they should be polynomial in n). Suppose that you have access to an algorithm $LP(G, s, t)$ to return the longest weighted path in G from s to t , (**c [1 point]**) explain how to solve the scheduling problem with the help of $LP()$ and finally (**d [1 point]**) explain in 4-5 lines how to implement $LP()$ so that the entire algorithm for task scheduling takes polynomial time, (**e [1 point]**) explain and derive the total complexity for the task scheduling problem using the above graph reduction approach.