Multi-threaded Program Analysis using "perf" Tool

Graduate Systems (CSE638) – PA02

1. Introduction

This project focuses on analyzing multithreaded programs using the Linux `perf` tool to assess performance, detect bottlenecks, and optimize execution. The study involves developing different types of workloads, implementing a producer-consumer pipeline, and evaluating performance metrics such as CPU usage, memory bandwidth, and context switches.

2. Objectives

- Develop four distinct multithreaded programs to simulate CPU-bound, memory-bound, I/O-bound, and mixed workloads.
- Implement a producer-consumer pipeline using synchronization primitives.
- Utilize `perf` to analyze execution efficiency and identify bottlenecks.
- Compare the scalability of different workload types by varying thread counts.

3. Implementation

3.1 Part A: Multithreaded Program Variants

Four separate programs were implemented using the pthread library:

1. **CPU-bound:** Performs intensive computations (e.g., matrix multiplication, prime number calculations).
2. **Memory-bound:** Focuses on operations requiring large memory allocations and accesses.
3. **I/O-bound:** Simulates disk I/O or network operations to assess read/write performance.
4. **Mixed workload:** Combines CPU, memory, and I/O tasks to simulate real-world applications.

Each program was executed with different thread configurations (2, 4, 8, 50, 100) to study the effect of thread scaling on performance.

4. Performance Analysis with `perf`

The following key metrics were analyzed for each program:

- **CPU Usage:** Measures processor load during execution.
- **Memory Bandwidth:** Tracks memory reads/writes per second.
- **Cache Behavior:** Monitors cache references and misses.
- **Context Switches:** Evaluates efficiency in handling multiple threads.
- **Execution Time & Throughput:** Determines overall performance.

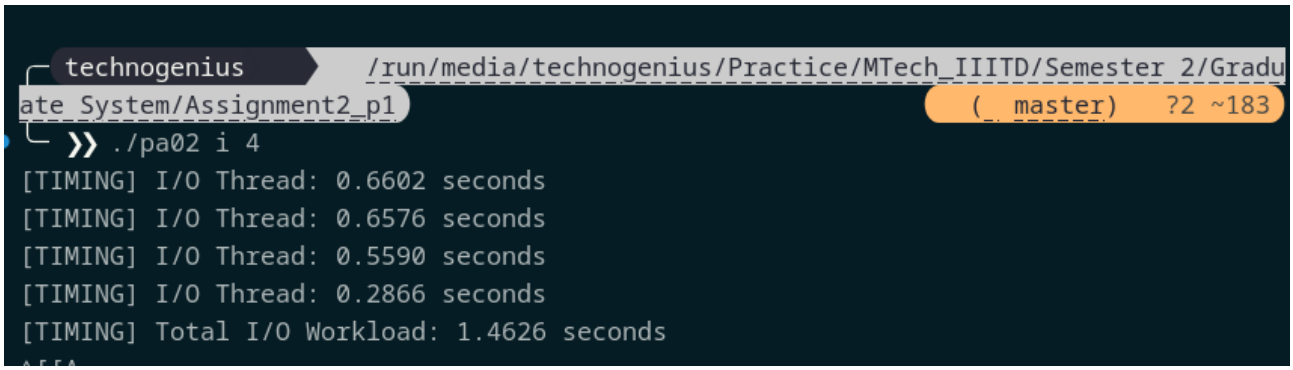

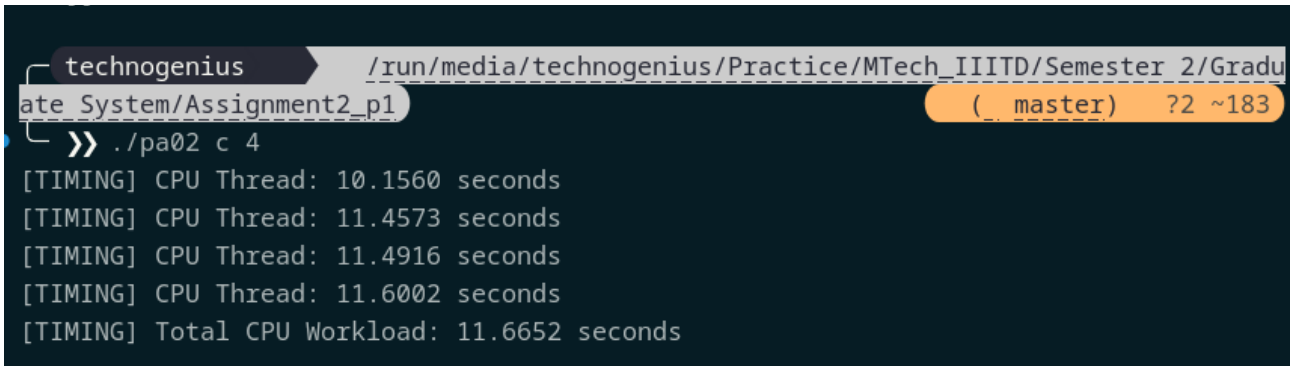

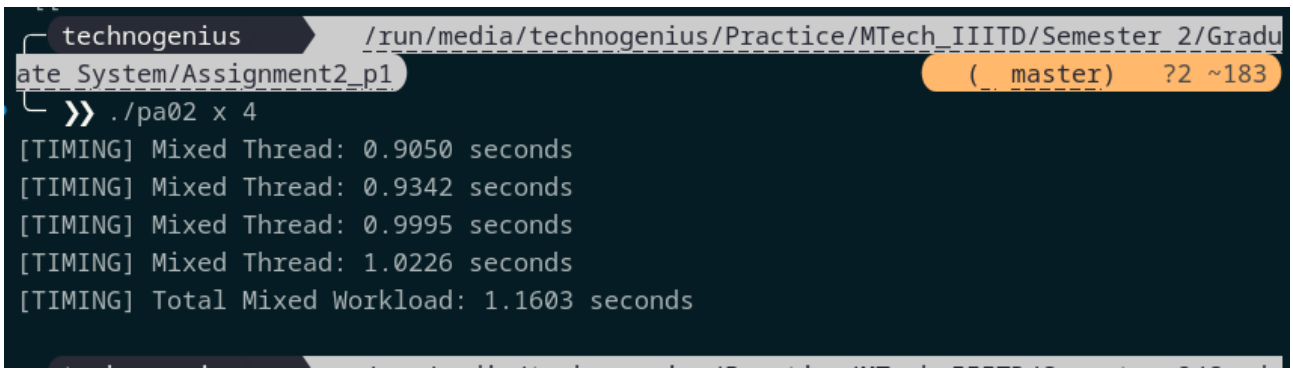

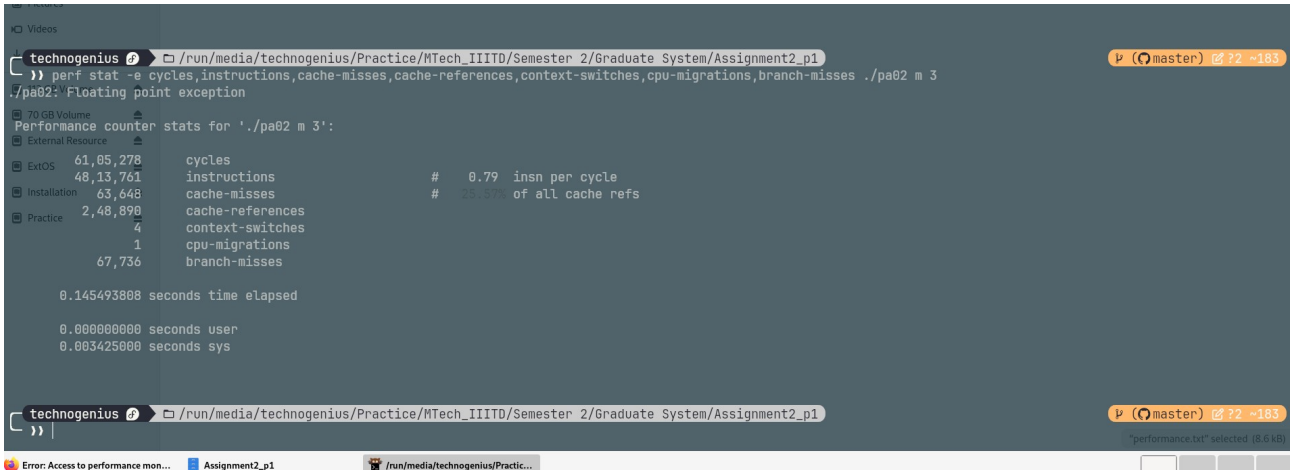

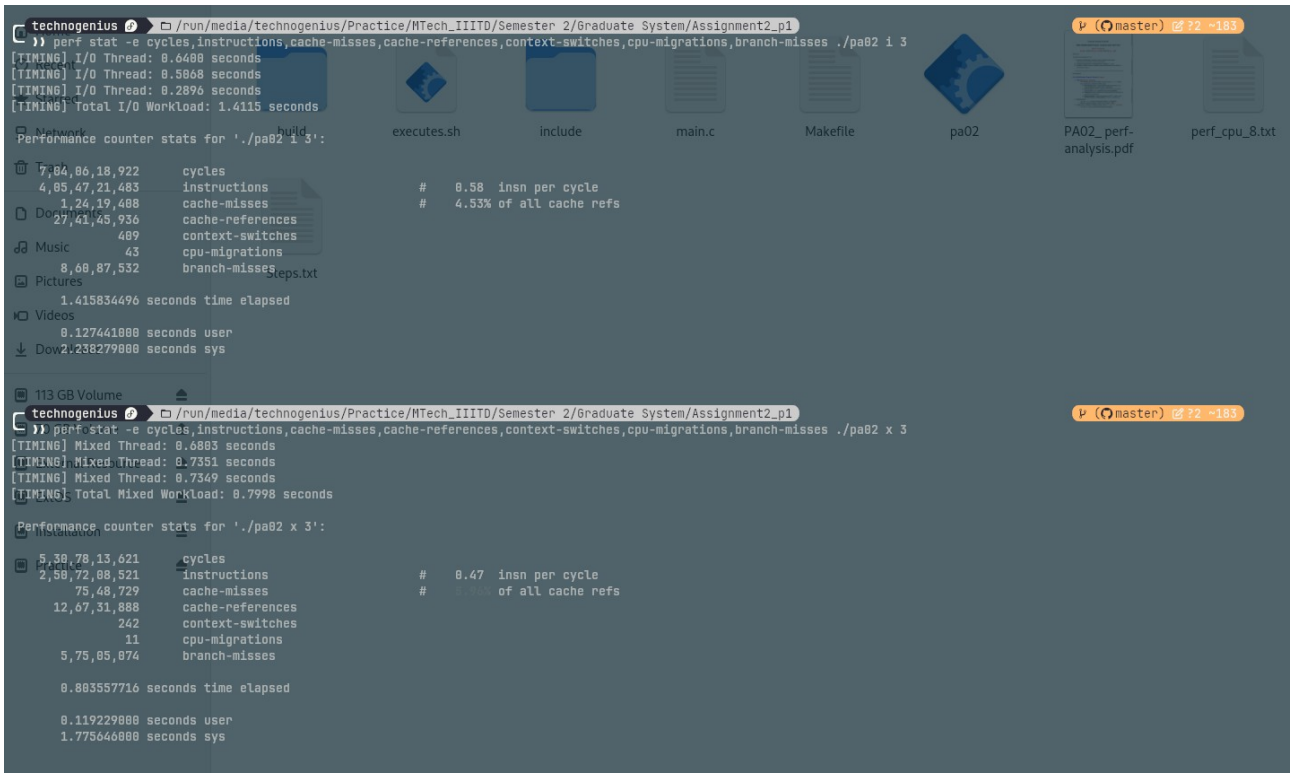

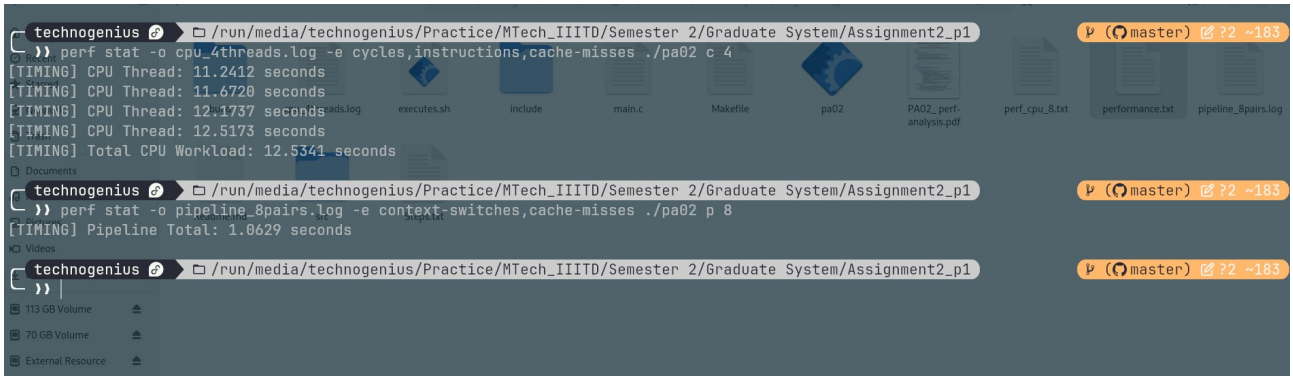

5. Scalability Testing

To assess scalability, each program was executed with increasing thread counts. Observations include:

- CPU-bound programs showed diminishing returns beyond a certain thread count.
- Memory-bound workloads suffered from increased cache misses with more threads.
- I/O-bound tasks exhibited bottlenecks due to disk read/write limitations.
- Mixed workloads demonstrated variable performance based on resource contention.

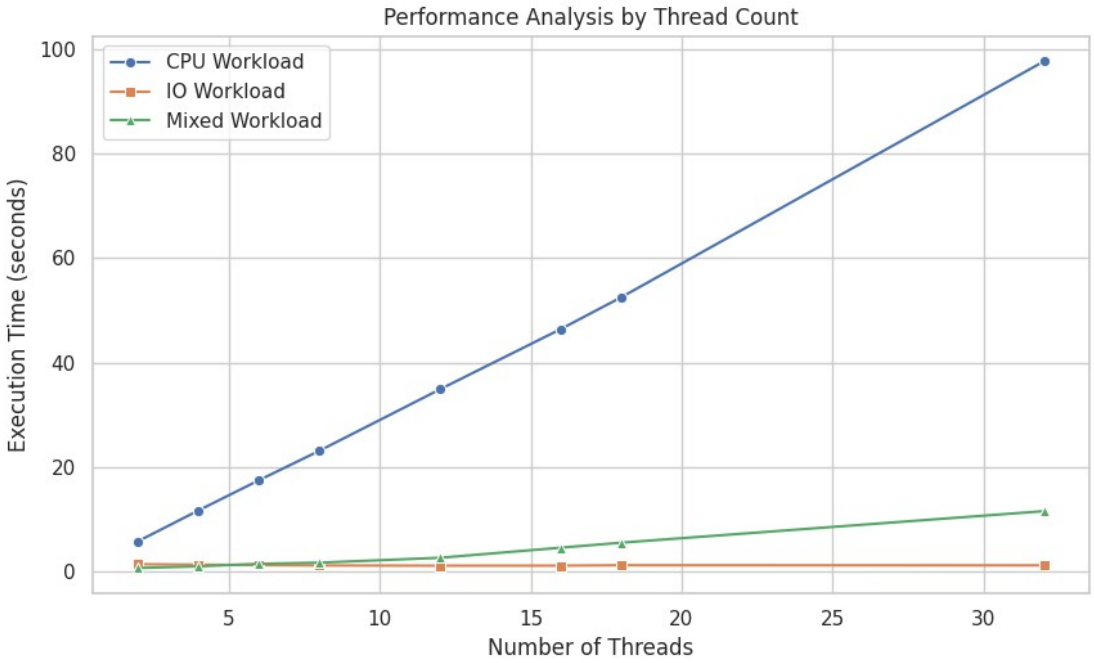

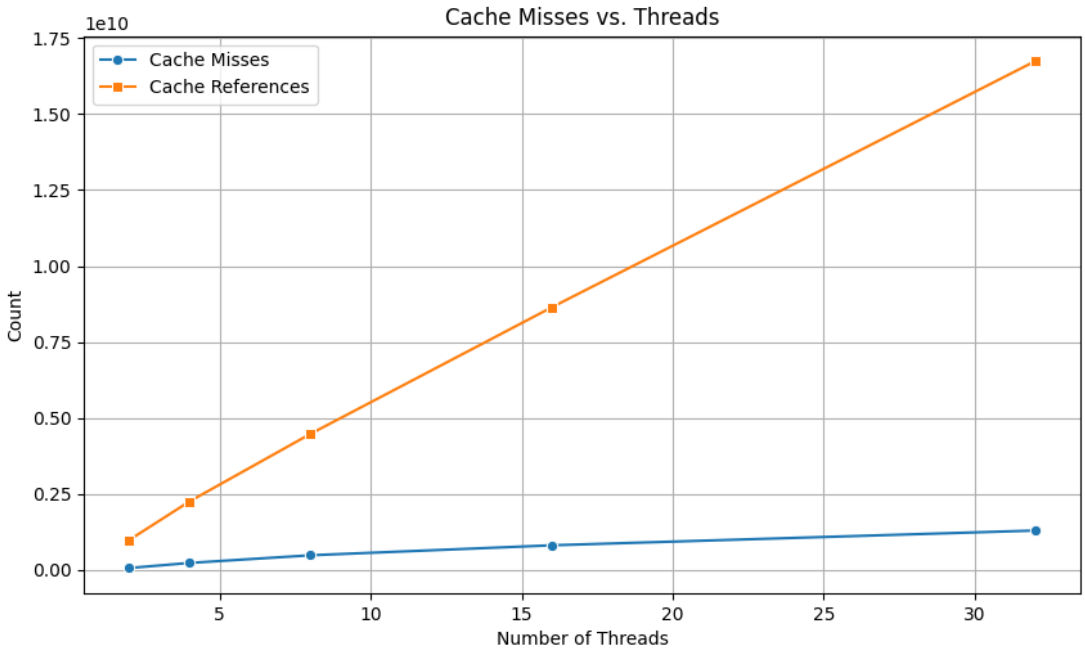

6. Observations and Bottlenecks

- Excessive thread creation led to overhead due to frequent context switching.
- Cache efficiency played a crucial role in memory-bound workloads.
- I/O operations became a limiting factor beyond a certain number of threads.
- The producer-consumer model improved efficiency but required careful synchronization tuning.

7. Conclusion

This project provided insights into multithreaded performance analysis using `perf`. The results highlighted the importance of workload-specific optimizations, proper thread management, and synchronization techniques for improving performance. Future improvements could include exploring alternative scheduling policies and optimizing memory access patterns.

8. References

- Linux `perf` Tool Documentation
- Multithreading with Pthreads
- Performance Optimization Techniques